

Graph Preconditioning for High-Dimensional Fixed Effects Regression

Alexander Fischer¹ and Kristof Schröder²

Draft: August 8, 2026

Abstract. Most software absorbs high-dimensional fixed effects with the Method of Alternating Projections (MAP). A MAP pass is cheap, but it updates one fixed-effect dimension at a time. In matched employer-employee data, sparse worker mobility can therefore make information travel slowly between groups of firms. We propose an alternative preconditioner for LSMR that uses the observed links between pairs of fixed effects. The worker-firm block of the fixed-effect cross product records worker-firm match counts; after a sign change, it is a weighted graph Laplacian. Sparse approximate Cholesky methods give inexpensive local corrections for this graph. We combine the pair corrections in an additive Schwarz preconditioner that can be reused across the outcome and covariates. In benchmarks, MAP and factor-pair LSMR are both fast on well-connected designs. As worker-firm connectivity weakens, MAP and diagonally preconditioned LSMR require substantially more time, while factor-pair LSMR remains fast. Its setup also becomes cheaper on the sparsest mobility graphs.

Keywords: high-dimensional fixed effects; alternating projections; preconditioning; matched employer-employee data; computational econometrics

JEL codes: C55; C63; C81; C87; J31

1. Introduction

Fixed effects appear in roughly half of the research published in leading economics and finance journals, according to Goldsmith-Pinkham (2026). Across applied economics, labor economists use worker and firm fixed effects to separate worker heterogeneity from firm wage premia; health economists study physician practice styles with individual, physician, and region fixed effects in mover designs; and education researchers use school, student, teacher, and student-teacher effects.

The Frisch-Waugh-Lovell (FWL) theorem reduces estimation to two operations: remove the variation explained by the fixed effects from the outcome and regressors, then run a low-dimensional regression on the residualized variables (Frisch and Waugh 1933; Lovell 1963). With high-dimensional fixed effects, software usually performs the first operation with the Method of Alternating Projections (MAP), also known as iterative demeaning or the “Zig-Zag” algorithm (Guimaraes and Portugal 2010; Gaure 2013). Packages including Stata’s `reghdfe` (Correia 2014; 2017), R’s `fixest` (Bergé et al. 2026), and Python’s `PyFixest` (PyFixest Developers 2026) use MAP or accelerated variants of it.

MAP updates one fixed-effect dimension at a time. In the worker-firm wage model of Abowd et al. (1999), extended here with year effects, one pass subtracts worker means, firm means, and year means in

¹trivago

²appliedAI Institute for Europe gGmbH

sequence. The firm update receives the residual left by the worker update, but it does not jointly solve the worker-firm problem.

That distinction matters when workers rarely move between groups of firms. Movers create the paths that separate worker effects from firm wage premia, while stayers add observations without connecting firms (Correia 2017; Jochmans and Weidner 2019). With few paths between groups, a change on one side of the graph reaches the other side only through successive MAP passes. The mobility pattern therefore affects identification, precision, and computation.

We follow Jochmans and Weidner (2019) and call the second-smallest eigenvalue of the normalized Laplacian the *Gap*. For factor pair (q, r) , let L_{qr} be the weighted graph Laplacian and Δ_{qr} its degree matrix. Then

$$S_{qr} = \Delta_{qr}^{-\frac{1}{2}} L_{qr} \Delta_{qr}^{-\frac{1}{2}}, \quad \text{Gap}_{qr} = \lambda_2(S_{qr}). \quad (1)$$

Edge weights equal the number of observed co-occurrences. A larger Gap means that the two fixed-effect dimensions are more strongly connected. We remove singleton levels before computing it. If the remaining graph is disconnected, we use the component containing the most observations and report its observation share. Unless stated otherwise, Gap refers to the worker-firm graph.

The fixed-effect cross product, or Gramian, contains this graph (Correia 2017). Its diagonal blocks are worker, firm, and year observation counts. Its off-diagonal blocks are match counts, including the worker-firm links that MAP handles only through successive residuals. Our preconditioner groups these blocks by factor pair. A sign change turns each pair block into a graph Laplacian, and sparse approximate Cholesky methods supply local corrections at low cost (Spielman and Teng 2014; Gao et al. 2025). Their weighted sum preconditions LSMR.³

The pair factorizations create a setup cost. They are unlikely to help when MAP already converges in a few passes, but they can save substantial iteration time on a weakly connected graph. Figure 1 reports total regression time as mobility and sorting change. The left column compares package defaults. The right column holds the package and achieved accuracy fixed within PyFixest, using tolerances calibrated in Section 7.

³FixedEffectModels.jl (FixedEffectModels.jl Developers 2026) also uses LSMR (Fong and Saunders 2011), with diagonal preconditioning based on fixed-effect counts. The factor-pair preconditioner also uses links between fixed-effect dimensions.




Figure 1: Median regression time against the worker-firm Gap $\lambda_{2,WF}$, on logarithmic axes. The Gap is computed after singleton removal on the connected component containing the most observations; larger values mean stronger connectivity. All panels use the same reversed horizontal scale, so connectivity weakens from left to right. Each simulated worker-firm-year panel has 1 million observations. The top row varies worker mobility, while the bottom row varies sorting among movers. The left column compares package defaults, including each package’s singleton rule. The right column compares four PyFixest configurations at tolerances that give comparable coefficient and residual errors in Section 7. A filled marker means all three planned fits returned estimates; a hollow marker means one or two did. Lines join observed medians and are not fitted trends. An arrow marks the median time at the iteration cap, which is a lower bound on completion time. Other failures are omitted.

All implementations finish quickly at high Gap values. MAP and unpreconditioned LSMR slow sharply as the Gap falls; factor-pair LSMR remains fast, and its runtime falls in the lowest-mobility designs because the local graphs become cheaper to factorize.

The next four sections introduce fixed-effect absorption, the AKM example, the Gramian, and the link between MAP and graph connectivity. Section 6 develops the factor-pair preconditioner. Sections 7 and 8 report the benchmarks and software interface, Section 9 concludes, and the appendix compares memory use.

2. Absorbing Fixed Effects⁴

We study the linear model

$$y = X\beta + D\alpha + \varepsilon, \quad (2)$$

where X contains the regressors of interest, D is the fixed-effect design matrix, and α collects the fixed-effect coefficients. Because D may have hundreds of thousands or millions of columns in high-dimensional applications, forming $[X \ D]$ or inverting its cross product can be computationally infeasible. The Frisch-Waugh-Lovell (FWL) theorem lets us compute $\hat{\beta}$ without either operation.

⁴Researchers employ several names for this operation: “absorbing fixed effects”, “demeaning”, “residualizing”, or applying the “within transformation”. We use these terms interchangeably throughout.

FWL removes the fixed-effect means from y and each column of X , then regresses the residualized outcome on the residualized covariates. Let M_D denote the operation that removes those means. In matrix terms, it is the weighted residual-making projection for the fixed-effect indicators, so $\tilde{y} = M_D y$ and $\tilde{X} = M_D X$. The second regression gives

$$\tilde{y} = M_D y, \quad \tilde{X} = M_D X, \quad \hat{\beta} = (\tilde{X}' W \tilde{X})^{-1} \tilde{X}' W \tilde{y}, \quad (3)$$

where W is a diagonal matrix of weights. With one covariate, this amounts to three regressions: y on D , x on D , and $\tilde{y}$ on $\tilde{x}$. FWL implies that the slope in the last regression equals the coefficient on x in the full regression of y on x and D .

To residualize the outcome and each covariate in X , we project them onto the column space of the fixed-effect dummy matrix D by weighted least squares. For a variable μ to be residualized - either y or one column of X - we solve

$$\hat{\alpha}_\mu = \arg \min_{\alpha} \| D\alpha - \mu \|_W^2, \quad (4)$$

and the residual is $\tilde{\mu} = \mu - D\hat{\alpha}_\mu$. The first-order condition for Equation 4,

$$D'W(D\hat{\alpha}_\mu - \mu) = 0, \quad G\hat{\alpha}_\mu = D'W\mu, \quad G = D'WD, \quad (5)$$

determines $\hat{\alpha}_\mu$. Its coefficient matrix G is the same for each FWL residualization; only the variable being residualized changes, so the computational cost depends on the structure of the Gramian G , the weighted cross-product matrix of the fixed-effect dummies. In the AKM worker-firm model, this structure has a concrete interpretation: worker moves link firms through shared workers and determine the off-diagonal blocks of G .

3. A Running Example: The AKM Model

Throughout the paper, we use the AKM model of Abowd et al. (1999) as the worker-firm example, in which workers observed at more than one firm provide the comparisons needed to distinguish persistent worker heterogeneity from firm wage premia. We write the regression as

$$y_{it} = \alpha_i + \psi_{J(i,t)} + \varphi_t + x'_{it}\beta + \varepsilon_{it}, \quad (6)$$

where α_i is a worker fixed effect, $\psi_{J(i,t)}$ is the fixed effect for the firm employing worker i at time t , and φ_t is a time fixed effect.

The AKM specification can be represented as a bipartite graph, with workers and firms as nodes and employment spells as edges. Worker moves induce a firm-to-firm graph by linking firms that employ the same worker; stayers - workers who never change employer - add observations to existing worker-firm links but do not bridge firms. Year effects enter as a third, low-dimensional factor on the same records, and Figure 2 contrasts a well-connected mobility graph with one that fragments under strong sorting.




Figure 2: Worker-firm graphs under high and low mobility. High mobility creates many paths between firms. With low mobility and strong sorting, only a few worker moves connect otherwise separate groups of firms.

A worker observed at one firm provides no comparison between worker productivity and the firm wage premium. Movers do: wages for the same worker at different firms help separate the two effects, especially when the move connects groups that otherwise share few workers. These comparisons are edges in the mobility graph and entries in the worker-firm cross-tabulation of the Gramian.

4. The Graph Structure of the Gramian

The block structure of the Gramian $G = D'WD$ represents the bipartite graph of worker and firm connections introduced in Section 3 (Correia 2017). Suppose that the columns of D are ordered as worker levels, firm levels, and year levels. Then

$$G = \begin{pmatrix} G_{WW} & C_{WF} & C_{WY} \\ C_{(WF)'} & G_{FF} & C_{FY} \\ C_{(WY)'} & C_{(FY)'} & G_{YY} \end{pmatrix}. \quad (7)$$

The **diagonal blocks** G_{WW} , G_{FF} , and G_{YY} contain weighted counts for workers, firms, and years. An observation belongs to one level of each factor, so these blocks are diagonal; solving them requires only division by group counts.

The **off-diagonal blocks** are cross-tabulations: the worker-firm block C_{WF} records how often worker i is observed at firm j , and the worker-year and firm-year blocks have analogous interpretations.

The small worker-firm panel below illustrates how its Gramian is constructed. We set all regression weights to one, so $W = I$.

Obs.	Worker	Firm	Year	y
1	W_1	F_1	Y_1	3.2
2	W_1	F_2	Y_2	4.1
3	W_2	F_1	Y_1	2.8
4	W_2	F_1	Y_2	3.9
5	W_3	F_2	Y_1	5.0
6	W_3	F_2	Y_2	4.5




Figure 3: Worker-firm graph for the example panel. Worker W_1 works at both firms and connects them. Worker W_2 works only at F_1 , and worker W_3 works only at F_2 .

Worker W_1 supplies the only link between F_1 and F_2 in Figure 3. Workers W_2 and W_3 stay at one firm. The diagonal blocks are count matrices. In this example, each worker is observed twice, each firm three times, and each year three times, so

$$G_{WW} = \begin{pmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 2 \end{pmatrix}, \quad G_{FF} = \begin{pmatrix} 3 & 0 \\ 0 & 3 \end{pmatrix}, \quad G_{YY} = \begin{pmatrix} 3 & 0 \\ 0 & 3 \end{pmatrix}. \quad (8)$$

The off-diagonal blocks are cross-tabulations between factors. The worker-firm block is

$$C_{WF} = \begin{pmatrix} 1 & 1 \\ 2 & 0 \\ 0 & 2 \end{pmatrix}. \quad (9)$$

The first row records the mover's two matches; the remaining rows record the two stayers. The other cross-tabulations are

$$C_{WY} = \begin{pmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{pmatrix}, \quad C_{FY} = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}. \quad (10)$$

Together, the count and cross-tabulation blocks form the full Gramian.

With column order $(W_1, W_2, W_3, F_1, F_2, Y_1, Y_2)$, the full Gramian is

$$G = \begin{pmatrix} 2 & 0 & 0 & 1 & 1 & 1 & 1 \\ 0 & 2 & 0 & 2 & 0 & 1 & 1 \\ 0 & 0 & 2 & 0 & 2 & 1 & 1 \\ 1 & 2 & 0 & 3 & 0 & 2 & 1 \\ 1 & 0 & 2 & 0 & 3 & 1 & 2 \\ 1 & 1 & 1 & 2 & 1 & 3 & 0 \\ 1 & 1 & 1 & 1 & 2 & 0 & 3 \end{pmatrix}. \quad (11)$$

The worker-firm submatrix contains worker and firm counts on its diagonal and match counts in C_{WF} . Changing the sign of C_{WF} makes the off-diagonal entries non-positive. Each row then sums to zero because its diagonal observation count equals the sum of its match counts. The result is the weighted graph Laplacian

$$L_{WF} = \begin{pmatrix} 2 & 0 & 0 & -1 & -1 \\ 0 & 2 & 0 & -2 & 0 \\ 0 & 0 & 2 & 0 & -2 \\ -1 & -2 & 0 & 3 & 0 \\ -1 & 0 & -2 & 0 & 3 \end{pmatrix}. \quad (12)$$

Every factor pair has a Laplacian of this form. Section 6 uses these matrices in the preconditioner; MAP uses only their diagonal count blocks in each update.

5. Alternating Projections and Graph Connectivity

The Method of Alternating Projections (MAP), also called iterative demeaning or the “zig-zag” algorithm, is the standard method for absorbing several fixed effects (Guimaraes and Portugal 2010; Gaure 2013). Implementations often accelerate the basic iteration (Berge 2018; Correia 2017); `fixest`, for example, uses Irons-Tuck extrapolation (Irons and Tuck 1969; Bergé et al. 2026).

In the worker-firm-year model, MAP starts with a partial residual and subtracts its mean within each worker. It next computes firm means from the updated residual, then year means. Repeating this pass eventually removes all three sets of means.

With $D = [D_W \ D_F \ D_Y]$, the FWL normal equations in Equation 5 are

$$\begin{pmatrix} G_{WW} & C_{WF} & C_{WY} \\ C_{(WF)'} & G_{FF} & C_{FY} \\ C_{(WY)'} & C_{(FY)'} & G_{YY} \end{pmatrix} \begin{pmatrix} \alpha_W \\ \alpha_F \\ \alpha_Y \end{pmatrix} = \begin{pmatrix} D_{W'} W \mu \\ D_{F'} W \mu \\ D_{Y'} W \mu \end{pmatrix}. \quad (13)$$

The worker update holds the current firm and year effects fixed. The worker row of this system can be written as

$$G_{WW} \alpha_W = D_{W'} W (\mu - D_F \alpha_F - D_Y \alpha_Y). \quad (14)$$

The diagonal entries of G_{WW} are workers’ total observation weights. Dividing the right-hand side by those counts gives the weighted worker means. Firm and year updates have the same form. One complete MAP pass applies these three inexpensive group-mean calculations in sequence.

The match-count blocks enter only through the changing partial residual. MAP never solves a worker-firm, worker-year, or firm-year block jointly.

With many overlapping employment histories, a worker update changes residuals at many firms, and the next firm update carries that information onward. Sparse mobility, strong sorting, or near nesting confines the changes to a narrower part of the graph. MAP may then need many passes even though each pass remains cheap.

The factor-pair Gap λ_{2,q^r} summarizes the strength of these links. A small Gap is a warning that information may move slowly across the pair graph. In a model with three or more fixed effects it remains a diagnostic for one pair, not a bound on the convergence rate of the full model.

6. The Factor-Pair Schwarz Preconditioner

6.1. Preconditioners

LSMR (Fong and Saunders 2011) allows each residual correction to use a preconditioner. Weak links, sparse mobility, or near nesting make some combinations of fixed effects much less well determined than others. Changing from MAP to LSMR does little to remove that imbalance; the preconditioner must rescale the system while preserving its least-squares solution.

The inverse Gramian G^{-1} shows the target. If $M^{-1} = G^{-1}$, LSMR sees the identity matrix,⁵

$$M^{-1}G = G^{-1}G = I. \quad (15)$$

Under this choice, one correction would suffice. Computing the full inverse costs as much as the original fixed-effect problem. A useful approximation must capture the poorly determined directions and save enough iterations to cover its construction and application costs.⁶

6.2. From the Block Inverse to the Diagonal Preconditioner

The block inverse identifies the information lost by diagonal scaling. For the worker-firm-year model,

$$G = \begin{pmatrix} G_{WW} & C_{WF} & C_{WY} \\ C_{(WF)'} & G_{FF} & C_{FY} \\ C_{(WY)'} & C_{(FY)'} & G_{YY} \end{pmatrix}, \quad (16)$$

has diagonal weighted-count blocks and off-diagonal match-count blocks. Focus on the worker-firm part,

$$G_{WF} = \begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix}. \quad (17)$$

Its inverse is

$$G_{WF}^{-1} = \begin{pmatrix} G_{WW}^{-1} + G_{WW}^{-1}C_{WF}S^{-1}C_{(WF)'}G_{WW}^{-1} & -G_{WW}^{-1}C_{WF}S^{-1} \\ -S^{-1}C_{(WF)'}G_{WW}^{-1} & S^{-1} \end{pmatrix}. \quad (18)$$

$$S = G_{FF} - C_{(WF)'}G_{WW}^{-1}C_{WF}. \quad (19)$$

The diagonal inverse divides by worker counts. Worker-firm match counts enter through the Schur complement S , a firm-side system that compares firms after accounting for their shared workers. This is the expensive part. Exact factorization can add many nonzero entries, with time and memory determined by the connected-component sizes and the amount of fill-in. Sparse mobility can make an unpreconditioned iteration slow while making the local factorization cheaper because there are fewer links. With three factors, the full inverse can contain all three pairwise cross-tabulations.

Diagonal preconditioning drops the Schur-complement corrections and keeps only the count inverses,

$$M_{\text{diag}}^{-1} = \text{diag}(G_{WW}^{-1}, G_{FF}^{-1}, G_{YY}^{-1}). \quad (20)$$

This is the preconditioner used in `FixedEffectModels.jl` (Fong and Saunders 2011; FixedEffectModels.jl Developers 2026). It corrects differences in group size. It cannot distinguish a firm linked to many employers from an equally large firm whose workers remain in one part of the mobility graph.

6.3. The Factor-Pair Schwarz Approximation

Additive Schwarz preconditioning adds corrections from smaller, overlapping problems (Xu 1992; Toselli and Widlund 2005). For the AKM model, the subproblems are worker-firm, worker-year, and firm-year.

⁵As in any model with several fixed effects, the level effects are pinned down only up to a normalization: we can add a constant to every worker effect and subtract it from every firm effect without changing the fitted values $D\alpha$, and likewise for years. The dummy-coded $G = D'WD$ and the smaller blocks inverted below are therefore not invertible until this ambiguity is removed - one normalization for the worker-firm pair, and a second once year effects are added, as in the Section 4 example. We adopt the standard normalization within each connected set and read every inverse below as that of the resulting system. The estimated effects depend on the normalization; the residualized outcome and regressors, which are all the regression uses, do not.

⁶LSMR never constructs $M^{-1}G$ or G explicitly. It multiplies vectors by D and D' and applies M^{-1} .

Each retains its match-count block, while the outer LSMR iteration handles the remaining three-way coupling.

The worker-firm local problem is

$$\begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix}. \quad (21)$$

Figure 4 compares this worker-firm pair block with the single diagonal block solved by a factor-level MAP update.




Figure 4: Matrices used by MAP and the factor-pair preconditioner in the example from Section 4. A MAP worker update (left) uses only the diagonal worker-count matrix $G_{WW} = \text{diag}(2, 2, 2)$. The factor-pair update (right) adds the firm-count matrix $G_{FF} = \text{diag}(3, 3)$ and the matrix of worker-firm observation counts C_{WF} . The dashed outline marks the worker and firm coefficients that are updated together. Its inverse retains the worker-firm links through the Schur complement. To place this correction in the three-factor problem, let R_{WF} select the worker and firm entries from the full coefficient vector $\alpha = [\alpha_W; \alpha_F; \alpha_Y]$; its transpose $R_{(WF)'}$ places the resulting correction back into the full vector. The diagonal matrix $\tilde{D}_{WF}$ weights levels that appear in more than one pair problem with partition-of-unity weights, so their corrections are not counted twice. The exact worker-firm contribution is

$$P_{WF}^{-1} = R_{(WF)'} \tilde{D}_{WF} \begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix}^{-1} \tilde{D}_{WF} R_{WF}. \quad (22)$$

The worker-year and firm-year terms have the same form. Each level belongs to two pair problems, so its weight in each is $\frac{1}{\sqrt{2}}$; the squared weights then sum to one. Adding the three terms gives the exact factor-pair Schwarz preconditioner,

$$P^{-1} = P_{WF}^{-1} + P_{WY}^{-1} + P_{FY}^{-1}. \quad (23)$$

The pair terms capture all three cross-tabulations separately. LSMR resolves the joint three-factor coupling and the error from splitting shared levels across pairs.

6.4. Approximating Pair Systems via Graph Laplacians

Direct inversion is practical only for small pair blocks. The Section 4 example reduces to a 4×4 system after normalization, but register data can contain hundreds of thousands of workers and firms. Large blocks require an approximation that uses their graph-Laplacian structure.

For a worker-firm pair, the local pair step solves the pair-Gramian system

$$\begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix} x = u, \quad (24)$$

where u contains the selected and weighted worker and firm entries of the LSMR vector. The match counts C_{WF} are non-negative, so this matrix is not yet a Laplacian. Let $T_{WF} = \text{diag}(I_W, -I_F)$ flip the firm signs. Multiplication on both sides gives

$$L_{WF} = T_{WF} \begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix} T_{WF} = \begin{pmatrix} G_{WW} & -C_{WF} \\ -C_{(WF)'} & G_{FF} \end{pmatrix}, \quad (25)$$

a weighted bipartite graph Laplacian with non-positive off-diagonals and zero row sums. Since $T_{WF}^2 = I$, the same sign change maps its inverse back to the pair Gramian,

$$\begin{pmatrix} G_{WW} & C_{WF} \\ C_{(WF)'} & G_{FF} \end{pmatrix}^{-1} = T_{WF} L_{WF}^{-1} T_{WF}, \quad (26)$$

where the inverses use the normalization from Section 6.1. The implementation returns the zero-mean solution within each connected component. This choice fixes the free constant but does not change the residualized variables.

The local correction need not be exact because LSMR refines the remaining error. We use sparse approximate Cholesky factorizations (Spielman and Teng 2014; Gao et al. 2025). Exact elimination can create fill-in: removing a worker links the firms that employed that worker, and later eliminations add more entries. A dense k -level factorization can require order k^3 operations and k^2 memory. Randomized approximation limits this growth, with cost close to the number of observed links up to logarithmic factors. Let A_{WF} denote the resulting approximate pair-Gramian inverse after the signs are changed back.

The worker-year and firm-year approximations are A_{WY} and A_{FY} . Substituting all three approximate inverses into the Schwarz sum gives the implemented preconditioner,

$$M^{-1} = \sum_{(q,r)} R_{(qr)'} \tilde{D}_{qr} A_{qr} \tilde{D}_{qr} R_{qr}. \quad (27)$$

The exact P^{-1} replaces the full three-factor inverse with exact pair inverses, while M^{-1} also approximates each pair inverse through A_{qr} . Both approximations change the update at each LSMR iteration but leave the original fixed-effect least-squares problem unchanged. LSMR therefore returns the same fitted residuals as an unpreconditioned LSMR run, up to the requested tolerance.

6.5. Implementation Strategy

Figure 5 summarizes the implementation. Each factor pair becomes a graph Laplacian after the sign change. Sparse approximate Cholesky supplies the local inverse, and partition-of-unity weights prevent shared levels from being counted twice in the Schwarz sum.

LSMR uses the sum to scale its updates (Fong and Saunders 2011; Arridge et al. 2014; Yang et al. 2024). The same preconditioner can residualize the outcome and every covariate when the observations, weights, and fixed-effect identifiers are unchanged. The within source implements this setup-and-apply sequence (within Developers 2026).




Figure 5: Construction of the factor-pair preconditioner. Each local problem combines two fixed-effect dimensions and retains their observed links. A sign change turns its matrix into a graph Laplacian, the standard matrix representation of a weighted graph. Sparse approximate Cholesky solves each local system while limiting the extra nonzero entries created during factorization, which saves memory and computation. The weighted sum of these local solutions is the preconditioner for LSMR.

7. Benchmarks

The benchmarks ask when fewer iterations repay the cost of building the factor-pair preconditioner. We vary mobility and sorting in controlled AKM panels, compare two public synthetic designs, and use eight real datasets from Correia (2017).

Elapsed times cover the package-level regression call, from model setup through coefficient estimation. Cross-package comparisons use the same input data but retain each package’s default treatment of singleton levels and separated observations. Experiments that isolate the preconditioner instead use one prepared sample. The appendix reports the additional memory used to store the pair structures.

Each table reports the factor-pair Gap after iterative singleton removal. For a disconnected graph, the reported value comes from the component with the most retained observations; the number in parentheses is its observation share. The Gap remains a pairwise diagnostic in models with three or more fixed effects.

The OLS package-runtime tables compare six configurations: PyFixest MAP, PyFixest LSMR with no, diagonal, or factor-pair preconditioning, R `fixest`, and `FixedEffectModels.jl`:

Configuration	Package	Algorithm
rust-map	PyFixest	Vanilla Rust MAP, without acceleration.
within-off	PyFixest	LSMR without preconditioning.
within-diagonal	PyFixest	LSMR with diagonal preconditioning.
within	PyFixest	LSMR with factor-pair Schwarz preconditioning.

Configuration	Package	Algorithm
fixest	R fixest	Accelerated MAP with Irons-Tuck and other optimizations (Bergé et al. 2026).
FEM.jl	FixedEffectModels.jl	Diagonally preconditioned LSMR (Fong and Saunders 2011; FixedEffectModels.jl Developers 2026).

Tables report medians over the stated number of repetitions. All runs use an Apple M4 Mac mini with 10 CPU cores and 16 GB of memory running macOS 15.3.1. We omit `reghdfe` (Correia 2014; 2017) because we lack a Stata license; its accelerated MAP algorithm is represented here by `fixest`.

7.1. Runtime Benchmarks

Package defaults are useful to users but do not impose equal accuracy because the implementations apply tolerances to different quantities. We report both package-default times and a comparison based on achieved coefficient and residual error.

The main OLS panels contain one million observations, one covariate, ten periods, and worker, firm, and year effects. One set of designs lowers the probability of moving. The other holds that probability at one but increasingly assigns moves within groups of similar firms.

7.1.1. Worker Mobility

The worker-firm Gap falls from 0.232 to 2.41×10^{-5} as mobility declines. All configurations finish within 3.60 seconds in the first two designs. PyFixest MAP then slows sharply or reaches its iteration cap, and unpreconditioned LSMR also reaches the cap. Factor-pair LSMR falls from 0.554 seconds in the first design to 0.373 seconds in the last because the local worker-firm graphs become cheaper to factorize.

7.1.2. Sorting Among Movers

Increasing ρ concentrates moves within groups even though every worker moves between periods. The Gap falls from 0.222 to 2.13×10^{-4} . PyFixest MAP rises from 0.411 seconds to 53.9 seconds before reaching its cap in the final design; unpreconditioned LSMR reaches its cap in the last four. Diagonal LSMR rises from 0.362 to 1.58 seconds. Factor-pair LSMR stays below one second and is faster than diagonal LSMR in the three designs with the smallest Gap.

7.1.3. When Preconditioner Setup Dominates

The public `fixest` benchmark suite complements the gradual connectivity changes in the AKM designs by contrasting a dense, well-connected design with a sparse, nearly nested design (Bergé et al. 2026). Both contain 10 million observations, one covariate, and worker, firm, and year fixed effects.

Simple and difficult fixest designs (10 million observations, three fixed effects).

Design	Gap λ_2 (share)	PyFixest MAP	PyFixest LSMR none	PyFixest LSMR diagonal	PyFixest LSMR factor-pair	fixest	FEM.jl
simple (well-connected)	0.622 (1.00)	2.35s	2.72s	2.16s	11.2s	2.52s	2.08s
difficult (near-nested)	8.35×10^{-8} (1.00)	326.2s	10.2s	capped (0/3)	4.63s	62.2s	27.0s

Note: Times are in seconds and cover each OLS regression from model setup through coefficient estimation. Each package uses its default settings. A preliminary run sets the number of measured runs: 20 if it takes less than 1 second, 7 if it takes 1 to 10 seconds, and 3 if it takes longer. When only k of n planned runs return an estimate, $t(\frac{k}{n})$ reports their median time t . A dash means that no timing is available. capped (0/ n) means that no run finishes before the iteration limit; failed (0/ n) marks a failure other than reaching that limit. Both designs have 10 million observations, one covariate, and worker, firm, and year fixed effects. Gap λ_2 is computed after iterative singleton removal on the component containing the most retained observations; its observation share appears in parentheses.

The dense design exposes the setup cost: five methods finish in 2.08 to 2.72 seconds, while factor-pair LSMR takes 11.2 seconds. Near nesting reverses the ranking. Factor-pair LSMR takes 4.63 seconds, compared with 10.2 seconds without preconditioning, 27.0 seconds for FEM.jl, 62.2 seconds for fixest, and 326.2 seconds for PyFixest MAP; diagonal LSMR reaches its default cap.

7.1.4. Iterations After Setup

The iteration-count diagnostic removes setup time and uses 100,000-observation versions of the same designs.

Iterations on the simple and difficult designs (100,000 observations).

Design	map (passes)	off	diagonal	additive
simple	14	37	16	14
difficult	8159	249	180	22

Note: The MAP column counts complete passes over all fixed-effect dimensions. The other columns count LSMR iterations: off uses no preconditioner, diagonal scales by fixed-effect group counts, and additive uses the factor-pair preconditioner. A MAP pass and an LSMR iteration perform different work, so their counts are not directly comparable. Each count is the median of three runs; each run removes the fixed effects from one outcome and one covariate.

Factor-pair LSMR needs 14 iterations on the simple design and 22 on the difficult one. The corresponding counts rise from 16 to 180 for diagonal LSMR and from 37 to 249 without preconditioning. The factor-pair iteration count is stable, but its setup still makes it unattractive for the simple one-shot regression.

7.1.5. Public Real-Data Benchmarks

The controlled designs isolate mobility, sorting, and near nesting, but empirical fixed-effect graphs also contain large differences in group size, thin links between dense groups, and many disconnected components. The public benchmark collection assembled by Correia (2017) contains eight such datasets and lets us check whether the controlled comparisons carry over to these less regular graphs.

Correia real-data benchmarks.

Dataset	Gap λ_2 (share)	PyFixest MAP	PyFixest LSMR factor-pair	fixest	FEM.jl
credit	0.227 (1.00)	0.193s	0.242s	0.154s	0.097s
soccer	0.666 (1.00)	0.026s	0.065s	0.014s	0.010s
enron	3.53×10^{-3} (0.99)	2.95s	0.402s	0.535s	0.317s
github	3.15×10^{-4} (0.32)	capped (0/3)	0.310s	1.43s	1.44s
patents	2.59×10^{-4} (0.95)	capped (0/3)	0.388s	1.05s	0.576s
workers	1.37×10^{-4} (0.63)	capped (0/3)	0.279s	2.00s	1.45s
schools	1.11×10^{-3} (1.00)	5.98s	0.266s	0.828s	0.487s
directors	2.56×10^{-4} (0.30)	capped (0/3)	0.327s	0.283s	0.462s

Note: Times are in seconds for OLS regressions using each package’s default settings, with three planned runs per cell. Each package applies its own rule for removing fixed-effect levels observed only once. If only k runs return an estimate, $t(\frac{k}{3})$ reports their median time t . capped (0/3) means that no run finishes before the 10,000-iteration limit. Gap λ_2 is computed after iterative singleton removal on the id1-id2 component containing the most retained observations; its observation share appears in parentheses.

On *credit* and *soccer*, PyFixest MAP takes 0.193 and 0.026 seconds, compared with 0.242 and 0.065 seconds for factor-pair LSMR, so the preconditioner’s setup cost does not pay on these easy datasets. On the more difficult graphs, factor-pair LSMR takes 0.402 seconds on *enron* and 0.266 seconds on *schools*, compared with 2.95 and 5.98 seconds for PyFixest MAP. On *github*, *patents*, *workers*, and *directors*, PyFixest MAP reaches its iteration cap while factor-pair LSMR finishes in 0.279 to 0.388 seconds. Accelerated MAP can still be competitive, however; *fixest* takes 0.283 seconds on *directors*, slightly less than factor-pair LSMR’s 0.327 seconds.

The Gap is a pairwise description, not a rule for choosing a solver. Runtime also depends on component size, factor dimensions, graph sparsity, acceleration, and the preconditioner’s setup cost.

7.2. Runtime and Achieved Precision

MAP implementations compare changes in residuals or fixed-effect coefficients, while LSMR uses its estimated least-squares residual and first-order-condition error (Fong and Saunders 2011). Their numerical tolerances are therefore not a common measure of accuracy. For example, within stops when either its relative residual or the scaled normal-equation error $\|A^T r\|_2 / \|A\|_F \|r\|_2$ meets the requested threshold, where $A = W^{\frac{1}{2}}D$ and $r = W^{\frac{1}{2}}(\mu - D\alpha)$.

Figure 6 places the methods on a common accuracy scale for mobility designs 1, 3, and 5. Every method receives the same singleton-pruned sample and a common 10,000-iteration cap. We time three fits at each package-specific tolerance and compare their coefficients and residuals with factor-pair LSMR at tolerance 10^{-14} .




Figure 6: Elapsed time and achieved accuracy on three AKM mobility designs. Each marker reports the median elapsed time in seconds among the three planned fits that return an estimate on the same sample. Before timing, we remove fixed-effect levels observed only once and use the remaining observations for every method. Each line connects results for one method across requested tolerances. The top row measures coefficient error as $\frac{\|\hat{\beta} - \hat{\beta}^*\|_2}{SE(\hat{\beta}^*)}$. The bottom row measures residual error as $\frac{\|r - r^*\|_2}{\|r^*\|_2}$. The reference coefficient $\hat{\beta}^*$ and residual r^* come from factor-pair LSMR at tolerance 10^{-14} . Error decreases from left to right; on each line, a circle marks the result obtained with the package's default tolerance. Every fit is limited to 10,000 iterations. Settings for which none of the three fits returns an estimate are omitted and identified in the annotations.

7.3. Amortizing the Preconditioner

Achieved accuracy does not remove the factor-pair setup cost. Later regressions can reuse the pair blocks and factorizations when the observations, weights, and fixed-effect identifiers are unchanged.

7.3.1. Setup Cost Across Connectivity

We separate construction from LSMR iterations on the one-million-observation mobility designs, for specifications with worker-firm and worker-firm-year effects.

Factor-pair setup and LSMR times as worker mobility varies.

Move probability δ	Gap λ_2 (share)	2 FE setup	2 FE LSMR	3 FE setup	3 FE LSMR
1	0.232 (1.00)	0.109s	0.291s	0.113s	0.209s
0.5	0.0392 (1.00)	0.115s	0.255s	0.118s	0.234s
0.05	1.84×10^{-3} (0.87)	0.032s	0.135s	0.034s	0.233s
0.01	5.35×10^{-4} (0.61)	0.027s	0.124s	0.031s	0.216s
0.005	5.50×10^{-5} (0.52)	0.025s	0.093s	0.031s	0.193s
0.001	2.41×10^{-5} (0.29)	0.028s	0.045s	0.032s	0.149s

Note: Times are in seconds. For each design and specification, we plan five runs with 1 million observations and an LSMR tolerance of 10^{-12} . The table gives the median setup time and time spent in LSMR among the runs that finish; if only k finish, $t(\frac{k}{5})$ reports their median time t . The two-fixed-effect specification absorbs worker and firm effects; the three-fixed-effect specification also absorbs year effects. Gap λ_2 is computed after iterative singleton removal on the component containing the most retained observations; its observation share appears in parentheses.

From the highest- to the lowest-mobility design, setup falls from 0.109 to 0.028 seconds with two fixed effects and from 0.113 to 0.032 seconds with three. Time spent in LSMR also falls. The low-mobility graph

has fewer cross-firm links to store and factorize, which explains the declining factor-pair runtime in Figure 1.

7.3.2. Ten Regressions on the Same Fixed Effects

Empirical projects often estimate several specifications on one sample and set of fixed effects. We time ten sequential regressions on both one-million-observation `fixest` designs, comparing diagonal preconditioning with factor-pair preconditioning that is rebuilt for each regression or retained for all ten.

Ten regressions with rebuilt and cached preconditioners.

Design	Policy	Setup	LSMR	Total	Speedup
simple	Diagonal	0.048s	0.553s	0.601s	1.0x
	Additive, rebuilt	2.80s	1.51s	4.29s	0.1x
	Additive, cached	0.281s	1.51s	1.79s	0.3x
difficult	Diagonal	0.049s	43.6s	43.6s	1.0x
	Additive, rebuilt	0.379s	1.20s	1.58s	27.5x
	Additive, cached	0.040s	1.26s	1.30s	33.6x

Note: For each policy, we add the setup and LSMR times across ten sequential regressions and take the median across three planned repetitions. If only k repetitions finish, $t(\frac{k}{3})$ reports their median time t , and the speedup is omitted. Times are in seconds. Each design has 1 million observations and worker, firm, and year fixed effects. Each regression removes the fixed effects from the common outcome and one of ten covariates at an LSMR tolerance of 10^{-12} . For complete cells, speedup divides the diagonal total time by the reported total time. `Additive, rebuilt` constructs a new factor-pair preconditioner for each regression; `Additive, cached` reuses one preconditioner for all ten.

Reuse lowers factor-pair time on the simple design from 4.29 to 1.79 seconds, but diagonal preconditioning still takes only 0.601 seconds. On the difficult design, the ranking is decisive: ten regressions take 43.6 seconds with diagonal preconditioning, 1.58 seconds when the factor-pair preconditioner is rebuilt, and 1.30 seconds when it is retained.

7.4. Poisson and Other GLMs

Poisson fixed-effect models repeat the absorption step at every iteratively reweighted least-squares (IRLS) iteration (Stammann 2018). The fixed-effect identifiers remain fixed while the weights and working response change, so residualization speed accumulates over the fit. This is the setting used by `pplhdfe` (Correia et al. 2020) and related estimators.

PyFixest currently retains the preconditioner from the first weighted demeaning call as the IRLS weights change. The benchmark reports that policy; it does not isolate reuse as the cause of the runtime difference. The comparison uses the one-million-observation simple and difficult `fixest` designs (Bergé et al. 2026). It includes R `fixest::fepois`, `GLFixedEffectModels.jl`, and PyFixest `fepois` with MAP or factor-pair LSMR. Each implementation receives the same 100-iteration IRLS limit and applies its default separation rule.

Poisson benchmarks (1 million observations, one covariate).

Design	FE	PyFixest MAP	fixest	GLFEM.jl	PyFixest LSMR factor-pair
simple (well-connected)	3	7.95s	4.63s	5.71s	9.74s
difficult (near-nested)	3	capped (0/3)	439.3s	129.8s	5.52s

Note: Times are in seconds for Poisson fixed-effect regressions with 1 million observations, one covariate, and worker, firm, and year fixed effects. `fixest` is R `fixest::fepois`; `rust-map` and `within` use PyFixest `fepois`; and `GLFEM.jl` is `GLFixedEffectModels.jl`. The `within` configuration reuses the first factor-pair preconditioner as the weights change across iteratively reweighted least squares steps. Each package applies its

default rule for removing separated observations. If only k of the three planned fits return an estimate, $t(\frac{k}{3})$ gives their median time t . capped (0/3) means that no fit finishes before either the package's inner demeaning limit or the common limit of 100 iteratively reweighted least squares steps; failed (0/3) marks another failure.

All four implementations finish the simple design in less than ten seconds. The difficult design separates them: PyFixest MAP reaches its iteration cap, GLFEM.jl takes 129.8 seconds, fixest takes 439.3 seconds, and factor-pair LSMR takes 5.52 seconds. Repeated absorption magnifies the cost of slow propagation across the worker-firm graph.

8. Software

The open-source within project provides the solver benchmarked in Section 7 (within Developers 2026), with a computational core written in Rust and APIs for Python and R.

Interface	Package	Registry	Install
Rust	within	crates.io	cargo add within
Python	within-py	PyPI	pip install within-py
R	withinr	py-econometrics R-universe	install.packages("withinr", repos = "https://py-econometrics.r-universe.dev")

Applied users can select the method from PyFixest without constructing the Gramian or its pair blocks (PyFixest Developers 2026):

```
import pyfixest as pf

fit = pf.feols(
    "y ~ x1 + x2 | worker + firm",
    data=df,
    demeaner=pf.LsmrDemeaner(),
)
```

The lower-level Python interface exposes the FWL operation. Here worker_code and firm_code are zero-based integer codes:

```
import numpy as np
from within import solve_batch

categories = np.asfortranarray(
    df[["worker_code", "firm_code"]].to_numpy(dtype=np.uint32)
)
variables = np.asfortranarray(df[["y", "x1", "x2"]].to_numpy())
res = solve_batch(categories, variables)
y_tilde, X_tilde = res.demeaned[:, 0], res.demeaned[:, 1:]
beta_hat = np.linalg.lstsq(X_tilde, y_tilde, rcond=None)[0]
```

solve_batch residualizes all three columns with one factor-pair preconditioner. The last line is the low-dimensional FWL regression from Section 2.

9. Conclusion

MAP's low cost per pass makes it the right default for many well-connected fixed-effect models. Its weakness appears when mobility is sparse, sorting is strong, or one effect is nearly nested in another. Successive mean updates then move information slowly through the graph.

Factor-pair preconditioning changes the update, not the regression. It uses observed co-occurrences to correct worker-firm and other pair blocks jointly, while LSMR preserves the original least-squares target. The benchmarks show that this setup cost is unnecessary on easy graphs but can replace hundreds of MAP passes or LSMR iterations on difficult ones. Sparse pair graphs can also be cheaper to factorize than dense ones.

The practical case is strongest when the factor-pair Gap is small or when the same fixed effects are absorbed repeatedly. Specification searches can reuse one preconditioner across outcomes and covariates, and IRLS estimators call the demeaning routine at every outer iteration. In these settings, graph structure is useful computational information, not only an identification diagnostic.

Appendix: Memory Use

MAP stores the current residuals and group sums. Factor-pair LSMR also retains match structures and local factorizations. We compare peak resident memory for the two Rust implementations inside PyFixest, using isolated processes so the surrounding regression code is held fixed.

Peak physical memory use (three fixed effects, one covariate).

Design	Gap λ_2 (share)	PyFixest MAP	PyFixest LSMR factor-pair
<i>100K observations</i>			
simple (well-connected)	0.622 (1.00)	354 MiB	411 MiB
difficult (near-nested)	8.35×10^{-6} (1.00)	357 MiB	383 MiB
<i>1M observations</i>			
simple (well-connected)	0.622 (1.00)	758 MiB	1,000 MiB
difficult (near-nested)	8.35×10^{-6} (1.00)	775 MiB	953 MiB

Note: For each regression, we record the largest amount of physical memory used by an isolated Python process, in MiB (2^{20} bytes). The table compares PyFixest OLS regressions using MAP or factor-pair LSMR. Each regression has one covariate and worker, firm, and year fixed effects. The results cover the simple and difficult designs at 100,000 and 1 million observations. Gap λ_2 is computed after iterative singleton removal on the component containing the most retained observations; its observation share appears in parentheses.

Factor-pair LSMR uses slightly more memory: 26-57 MiB at 100K observations and 178-242 MiB at 1M. The additional storage holds the pairwise co-occurrences, partition weights, and approximate Cholesky factors.

References

- Abowd, John M., Francis Kramarz, and David N. Margolis. 1999. "High Wage Workers and High Wage Firms." *Econometrica* 67 (2): 251–333. <https://doi.org/10.1111/1468-0262.00020>.
- Arridge, Simon R., Marta M. Betcke, and Lauri Harhanen. 2014. "Iterated Preconditioned LSQR Method for Inverse Problems on Unstructured Grids." *Inverse Problems* 30 (7): 75009. <https://doi.org/10.1088/0266-5611/30/7/075009>.
- Berge, Laurent. 2018. *Efficient Estimation of Maximum Likelihood Models with Multiple Fixed-Effects: The R Package Fenmlm*. Technical Report Nos. 2018–13.
- Bergé, Laurent R., Kyle Butts, and Grant McDermott. 2026. "Fixest: A Fast and Feature-Rich Framework for Econometric Estimations in R." <https://doi.org/10.48550/arXiv.2601.21749>.
- Correia, Sergio. 2014. "Reghdfe: Stata Module to Perform Linear or Instrumental-Variable Regression Absorbing Any Number of High-Dimensional Fixed Effects." <https://github.com/sergiocorreia/reghdfe>.
- Correia, Sergio. 2017. *Linear Models with High-Dimensional Fixed Effects: An Efficient and Feasible Estimator*. Technical report. <https://hdl.handle.net/10.2139/ssrn.3129010>.
- Correia, Sergio, Paulo Guimarães, and Tom Zylkin. 2020. "Fast Poisson Estimation with High-Dimensional Fixed Effects." *The Stata Journal* 20 (1): 95–115. <https://doi.org/10.1177/1536867X20909691>.
- FixedEffectModels.jl Developers. 2026. "Fixedeffectmodels.jl: Fast Estimation of Linear Models with High Dimensional Categorical Variables." <https://github.com/FixedEffects/FixedEffectModels.jl>.
- Fong, David Chin-Lung, and Michael Saunders. 2011. "LSMR: An Iterative Algorithm for Sparse Least-Squares Problems." *SIAM Journal on Scientific Computing* 33 (5): 2950–71. <https://doi.org/10.1137/10079687X>.
- Frisch, Ragnar, and Frederick V. Waugh. 1933. "Partial Time Regressions as Compared with Individual Trends." *Econometrica* 1 (4): 387–401. <https://doi.org/10.2307/1907330>.
- Gao, Yuan, Rasmus Kyng, and Daniel A. Spielman. 2025. "AC(k): Robust Solution of Laplacian Equations by Randomized Approximate Cholesky Factorization." *SIAM Journal on Scientific Computing*. <https://rasmuskyng.com/papers/GKS25.pdf>.
- Gaure, Simen. 2013. "OLS with Multiple High Dimensional Category Variables." *Computational Statistics & Data Analysis* 66 : 8–18. <https://doi.org/10.1016/j.csda.2013.03.024>.
- Goldsmith-Pinkham, Paul. 2026. *Tracking the Credibility Revolution across Fields*. Technical report.
- Guimaraes, Paulo, and Pedro Portugal. 2010. "A Simple Feasible Procedure to Fit Models with High-Dimensional Fixed Effects." *The Stata Journal* 10 (4): 628–49. <https://doi.org/10.1177/1536867X1101000406>.
- Irons, Bruce M., and Richard C. Tuck. 1969. "A Version of the Aitken Accelerator for Computer Iteration." *International Journal for Numerical Methods in Engineering* 1 (3): 275–77. <https://doi.org/10.1002/nme.1620010306>.
- Jochmans, Koen, and Martin Weidner. 2019. "Fixed-Effect Regressions on Network Data." *Econometrica* 87 (5): 1543–60. <https://doi.org/10.3982/ECTA14605>.

- Lovell, Michael C. 1963. "Seasonal Adjustment of Economic Time Series and Multiple Regression Analysis." *Journal of the American Statistical Association* 58 (304): 993–1010. <https://doi.org/10.1080/01621459.1963.10480682>.
- PyFixest Developers. 2026. "Pyfixest: Fast High-Dimensional Fixed Effects Regression in Python." <https://github.com/py-econometrics/pyfixest>.
- Spielman, Daniel A., and Shang-Hua Teng. 2014. "Nearly Linear Time Algorithms for Preconditioning and Solving Symmetric, Diagonally Dominant Linear Systems." *SIAM Journal on Matrix Analysis and Applications* 35 (3): 835–85. <https://doi.org/10.1137/090771430>.
- Stammann, Amrei. 2018. "Fast and Feasible Estimation of Generalized Linear Models with High-Dimensional K-Way Fixed Effects." <https://doi.org/10.48550/arXiv.1707.01815>.
- Toselli, Andrea, and Olof Widlund. 2005. *Domain Decomposition Methods: Algorithms and Theory*. Springer. <https://doi.org/10.1007/b137868>.
- within Developers. 2026. "Within: High-Performance Fixed-Effects Solver for Econometric Panel Data." <https://github.com/py-econometrics/within>.
- Xu, Jinchao. 1992. "Iterative Methods by Space Decomposition and Subspace Correction." *SIAM Review* 34 (4): 581–613. <https://doi.org/10.1137/1034116>.
- Yang, Mei, Gul Karaduman, and Ren-Cang Li. 2024. "Flexible Modified LSMR for Least Squares Problems."